SGD with Large Step Sizes Learns Sparse Features
Seminar Optimization

Popović Milutin

Supervisor: Radu Ioan Bot

31. October 2023

- Objective is to minimize functions of the form

$$f(x) = \frac{1}{n} \sum_{i=1}^{n} f_i(x)$$

■ Training Data:

$$\{(x_1, y_1), \ldots, (x_n, y_n)\} \in \mathbb{R}^d \times \mathscr{Y}$$

■ In large-scale ML: large dimension d
and large number of training data n.

- Training Data:

$$\{(x_1, y_1), \ldots, (x_n, y_n)\} \in \mathbb{R}^d \times \mathcal{Y}$$

- In large-scale ML: large dimension d
  and large number of training data n.

SGD (Stochastic gradient descent)

Classical examples of fitting the data via minimizing:

- Least Squares

$$\frac{1}{n}\|Ax - b\|_2^2 = \frac{1}{n}\sum_{i=1}^{n}(a_i^T x - b_i)^2$$

- Support Vector Machine (SVM):

$$\frac{1}{2}\|x\|_2^2 + \frac{C}{n}\sum_{i}^{n} max(0, 1 - y_i(x^T a_i + b))$$

- Deep Neural Nets

$$\frac{1}{n}\sum_{i}^{n} loss(y_i, DNN(x; a_i))$$

SGD (Stochastic gradient descent)

Classical examples of fitting the data via minimizing:

- Least Squares

$$\frac{1}{n}\|Ax - b\|_2^2 = \frac{1}{n}\sum_{i=1}^{n}(a_i^T x - b_i)^2$$

- Support Vector Machine (SVM):

$$\frac{1}{2}\|x\|_2^2 + \frac{C}{n}\sum_{i}^{n} max(0, 1 - y_i(x^T a_i + b))$$

- Deep Neural Nets

$$\frac{1}{n}\sum_{i}^{n} loss(y_i, DNN(x; a_i))$$

SGD (Stochastic gradient descent)

Classical examples of fitting the data via minimizing:

- Least Squares

$$\frac{1}{n}\|Ax - b\|_2^2 = \frac{1}{n}\sum_{i=1}^{n}(a_i^T x - b_i)^2$$

- Support Vector Machine (SVM):

$$\frac{1}{2}\|x\|_2^2 + \frac{C}{n}\sum_{i}^{n} max(0, 1 - y_i(x^T a_i + b))$$

- Deep Neural Nets

$$\frac{1}{n}\sum_{i}^{n} loss(y_i, DNN(x; a_i))$$

SGD (Stochastic gradient descent)

Classical examples of fitting the data via minimizing:

- Least Squares

$$\frac{1}{n}\|Ax - b\|_2^2 = \frac{1}{n}\sum_{i=1}^{n}(a_i^T x - b_i)^2$$

- Support Vector Machine (SVM):

$$\frac{1}{2}\|x\|_2^2 + \frac{C}{n}\sum_{i}^{n} max(0, 1 - y_i(x^T a_i + b))$$

- Deep Neural Nets

$$\frac{1}{n}\sum_{i}^{n} \text{loss}(y_i, DNN(x; a_i))$$

- Common pattern:

$$f(x) = \frac{1}{n} \sum_{i=1}^{n} f_i(x)$$

GD vs SGD

- GD would compute the gradient of every $f_i(x)$ to update the next iterate
- SGD picks a pseudorandom $i(r) \in \{1, 2, \ldots, n\}$
- then uses only $\nabla f_{i(r)}(x_k)$ as its descent direction

$$x^{k+1} = x^k - t_k \nabla f_{i(r)}(x^k)$$

- Key property : $\mathbb{E}[\nabla f_{i(r)}(x)] = \nabla f(x)$
- $\nabla f_{i(r)}(x)$ is an unbiased estimator !

GD vs SGD

- GD would compute the gradient of every $f_i(x)$ to update the next iterate
- SGD picks a pseudorandom $i(r) \in \{1, 2, \ldots, n\}$
- then uses only $\nabla f_{i(r)}(x_k)$ as its descent direction

$$x^{k+1} = x^k - t_k \nabla f_{i(r)}(x^k)$$

- Key property : $\mathbb{E}[\nabla f_{i(r)}(x)] = \nabla f(x)$
- $\nabla f_{i(r)}(x)$ is an unbiased estimator !

GD vs SGD

- GD would compute the gradient of every $f_i(x)$ to update the next iterate
- SGD picks a pseudorandom $i(r) \in \{1, 2, \ldots, n\}$
- then uses only $\nabla f_{i(r)}(x_k)$ as its descent direction

$$x^{k+1} = x^k - t_k \nabla f_{i(r)}(x^k)$$

- Key property : $\mathbb{E}[\nabla f_{i(r)}(x)] = \nabla f(x)$
- $\nabla f_{i(r)}(x)$ is an unbiased estimator !

GD vs SGD

- GD would compute the gradient of every $f_i(x)$ to update the next iterate
- SGD picks a pseudorandom $i(r) \in \{1, 2, \ldots, n\}$
- then uses only $\nabla f_{i(r)}(x_k)$ as its descent direction

$$x^{k+1} = x^k - t_k \nabla f_{i(r)}(x^k)$$

- Key property : $\mathbb{E}[\nabla f_{i(r)}(x)] = \nabla f(x)$
- $\nabla f_{i(r)}(x)$ is an unbiased estimator !

GD vs SGD

- GD would compute the gradient of every $f_i(x)$ to update the next iterate
- SGD picks a pseudorandom $i(r) \in \{1, 2, \ldots, n\}$
- then uses only $\nabla f_{i(r)}(x_k)$ as its descent direction

$$x^{k+1} = x^k - t_k \nabla f_{i(r)}(x^k)$$

- Key property : $\mathbb{E}[\nabla f_{i(r)}(x)] = \nabla f(x)$
- $\nabla f_{i(r)}(x)$ is an unbiased estimator !

Large stepsizes induce Sparse Features

- large step sizes -> *loss stabilization*
- the longer the larger step size is used the better the sparse representation

Figure: ResNet-18 (Residual Network with 18 layers) trained on CIFAR-10 (60k 32x32 images)

[1]

Large stepsizes induce Sparse Features

- large step sizes -> *loss stabilization*
- the longer the larger step size is used the better the sparse representation

Figure: ResNet-18 (Residual Network with 18 layers) trained on CIFAR-10 (60k 32x32 images)

[1]

Bibliography

[1] Maksym Andriushchenko et al. *SGD with Large Step Sizes Learns Sparse Features*. 2023. arXiv: 2210.05337 [cs.LG].

To be continued...

Thank You!